

Model Selection, Checking & State Decoding

hmezer

2025-07-31

We consider a time series giving the step lengths of the female elephant Habiba in the package ‘moveHMM’.

```
install.packages("moveHMM")
```

```
## Installing package into '/cloud/lib/x86_64-pc-linux-gnu-library/4.5'  
## (as 'lib' is unspecified)
```

```
library(moveHMM)
```

```
rawdata <- read.table("http://www.rolandlangrock.com/HMMs/elephant_rawdata.txt",  
header=T)  
head(rawdata)
```


```
##           x           y           timestamp  
## 21973 340611.9 62801.22 2014-01-23 11:30:42.000  
## 21974 340624.2 62828.47 2014-01-23 11:45:11.000  
## 21975 340633.1 62857.41 2014-01-23 12:00:13.000  
## 21976 340659.8 62860.71 2014-01-23 12:15:13.000  
## 21977 340852.4 62804.10 2014-01-23 12:30:11.000  
## 21978 340988.1 62742.88 2014-01-23 12:45:13.000
```

```
data <- prepData(rawdata, type="UTM", coordNames=c("x", "y"))  
head(data)
```

```
##      ID      step      angle      x      y      timestamp  
## 1 Animal1 29.87805      NA 340611.9 62801.22 2014-01-23 11:30:42.000  
## 2 Animal1 30.28453 0.1237818 340624.2 62828.47 2014-01-23 11:45:11.000  
## 3 Animal1 26.91687 -1.1492332 340633.1 62857.41 2014-01-23 12:00:13.000  
## 4 Animal1 200.69303 -0.4088491 340659.8 62860.71 2014-01-23 12:15:13.000  
## 5 Animal1 148.93990 -0.1376260 340852.4 62804.10 2014-01-23 12:30:11.000  
## 6 Animal1 15.82305 -0.5497463 340988.1 62742.88 2014-01-23 12:45:13.000
```


To determine the number of states and candidate parameters for the initial distribution, we consult the histogram:

```
plot(data$step)
```



```
hist(data$step, breaks=100)
```

Histogram of data\$step



We fit a 2-state model:

```
stepMean2 <- c(25, 100) # mean vector of gamma distribution (step lengths)
stepSD2 <- c(10, 30) # SD vector of gamma distribution (step lengths)
stepPar2 <- c(stepMean2, stepSD2)
mod2 <- fitHMM(data,
```

```

        nbStates=2,
        stepPar0=stepPar2,
        verbose=1,
        angleDist="none",
        stationary=T)

## iteration = 0
## Step:
## [1] 0 0 0 0 0 0
## Parameter:
## [1] 3.218876 4.605170 2.302585 3.401197 -1.500000 -1.500000
## Function Value
## [1] 11160.99
## Gradient:
## [1] 3082.62143 566.72336 -2444.69331 -6131.00806 17.99539 -13.91029

## Warning in nlm(f = nLogLike, p = wpar, nbStates = nbStates, bounds = bounds, :
## NA/NaN replaced by maximum positive value
## Warning in nlm(f = nLogLike, p = wpar, nbStates = nbStates, bounds = bounds, :
## NA/NaN replaced by maximum positive value

## iteration = 33
## Parameter:
## [1] 3.365238 5.229511 3.193954 4.942918 -2.641668 -2.117366
## Function Value
## [1] 7922.079
## Gradient:
## [1] -4.392401e-04 -3.961611e-04 5.323197e-04 1.354384e-04 3.228045e-05
## [6] 5.045170e-05
##
## Relative gradient close to zero.
## Current iterate is probably solution.

```

We fit a 3-state model:

```

stepMean3 <- c(5, 50, 150) # mean vector of gamma distribution (step lengths)
stepSD3 <- c(5, 10, 50) # SD vector of gamma distribution (step lengths)
stepPar3 <- c(stepMean3, stepSD3)
mod3 <- fitHMM(data,
               nbStates=3,
               stepPar0=stepPar3,
               verbose=1,
               angleDist="none",
               stationary=T)

## iteration = 0
## Step:
## [1] 0 0 0 0 0 0 0 0 0 0 0 0
## Parameter:
## [1] 1.609438 3.912023 5.010635 1.609438 2.302585 3.912023 -1.500000
## [8] -1.500000 -1.500000 -1.500000 -1.500000 -1.500000
## Function Value
## [1] 9294.366
## Gradient:
## [1] -568.508294 1564.336528 220.587067 -306.597448 -677.672830
## [6] -1981.787414 -41.382820 32.458988 -70.684296 -5.418658

```

```
## [11] 40.281062 -2.595033
## Warning in nlm(f = nLogLike, p = wpar, nbStates = nbStates, bounds = bounds, :
## NA/NaN replaced by maximum positive value
## iteration = 88
## Parameter:
## [1] 3.188337 4.640428 5.788146 2.979941 4.146594 5.215611 -2.322311
## [8] -5.269481 -1.629989 -2.175968 -4.255974 -1.192184
## Function Value
## [1] 7842.911
## Gradient:
## [1] -5.539686e-04 -6.148323e-04 -2.603653e-04 1.629798e-04 8.633040e-04
## [6] 8.457397e-05 -3.575613e-04 3.676309e-05 1.422839e-04 -9.195395e-06
## [11] 4.060269e-06 -2.113181e-04
##
## Relative gradient close to zero.
## Current iterate is probably solution.
```

Finally, we fit a 4-state HMM:

```
stepMean4 <- c(5, 30, 100, 300) # mean vector of gamma distribution (step lengths)
stepSD4 <- c(5, 10, 25, 150) # SD vector of gamma distribution (step lengths)
stepPar4 <- c(stepMean4, stepSD4)
mod4 <- fitHMM(data,
               nbStates=4,
               stepPar0=stepPar4,
               verbose=1,
               angleDist="none",
               stationary=T)
```

```
## iteration = 0
## Step:
## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
## Parameter:
## [1] 1.609438 3.401197 4.605170 5.703782 1.609438 2.302585 3.218876
## [8] 5.010635 -1.500000 -1.500000 -1.500000 -1.500000 -1.500000 -1.500000
## [15] -1.500000 -1.500000 -1.500000 -1.500000 -1.500000 -1.500000
## Function Value
## [1] 8228.915
## Gradient:
## [1] -362.459043 139.055795 584.544515 216.753221 122.660725 -466.291334
## [7] -240.968567 -105.110637 -39.719092 16.976745 32.444694 -16.922892
## [13] 5.100000 39.160475 19.546411 -36.022514 -1.866028 34.350085
## [19] 6.854785 -7.432995
## Warning in nlm(f = nLogLike, p = wpar, nbStates = nbStates, bounds = bounds, :
## NA/NaN replaced by maximum positive value
## iteration = 136
## Parameter:
## [1] 2.312195 3.612622 4.861916 5.905087 1.956030 3.114501
## [7] 4.210841 5.271723 -1.104223 -2.603545 -4.824480 -1.654802
## [13] -2.087830 -4.423115 -5.911706 -1.429120 -2.186579 -15.376580
## [19] -2.398733 -1.236545
## Function Value
## [1] 7786.559
```

```
## Gradient:
## [1] -1.142673e-03 -1.220004e-03  4.169681e-04  7.266610e-04  1.697605e-03
## [6] -3.460430e-04  2.481712e-04 -3.800687e-04  2.380353e-04 -7.929777e-05
## [11]  3.280189e-05  0.000000e+00 -2.121456e-04 -1.583298e-05  2.769235e-06
## [16]  1.406448e-04  1.580588e-05  1.892738e-06 -7.583126e-06 -1.676970e-04
##
## Relative gradient close to zero.
## Current iterate is probably solution.
```

Which model is the best fit among these three? We can utilize the information criteria to guide us:

```
# AIC
aic <- numeric(3)
aic[1] <- 2*mod2$mod$minimum + 2*length(mod2$mod$estimate)
aic[2] <- 2*mod3$mod$minimum + 2*length(mod3$mod$estimate)
aic[3] <- 2*mod4$mod$minimum + 2*length(mod4$mod$estimate)
aic
```

```
## [1] 15856.16 15709.82 15613.12
```

```
which.min(aic)
```

```
## [1] 3
```



```
# BIC
bic <- numeric(3)
T <- dim(data)[1]
bic[1] <- 2*mod2$mod$minimum + log(T)*length(mod2$mod$estimate)
bic[2] <- 2*mod3$mod$minimum + log(T)*length(mod3$mod$estimate)
bic[3] <- 2*mod4$mod$minimum + log(T)*length(mod4$mod$estimate)
bic
```

```
## [1] 15888.12 15773.76 15719.67
```

```
which.min(bic)
```

```
## [1] 3
```

```
par(mfrow=c(1,2))
plot(2:4,aic,type='b',main='AIC values',xlab='number of states N',ylab='AIC')
plot(2:4,bic,type='b',main='BIC values',xlab='number of states N',ylab='BIC')
```

AIC values**BIC values**

```
mod2$mle
```

```
## $stepPar
##      state 1  state 2
## mean 28.94037 186.7015
## sd   24.38466 140.1787
##
## $anglePar
## NULL
##
## $beta
##           1 -> 2    2 -> 1
## intercept -2.641668 -2.117366
##
## $delta
## [1] 0.6176254 0.3823746
##
## $gamma
##           [,1]      [,2]
## [1,] 0.9334956 0.06650444
## [2,] 0.1074204 0.89257962
```

```
mod3$mle
```

```
## $stepPar
##      state 1  state 2  state 3
## mean 24.24808 103.58869 326.4074
## sd   19.68666  63.21832 184.1243
##
## $anglePar
## NULL
```

```

##
## $beta
##           1 -> 2    1 -> 3    2 -> 1    2 -> 3    3 -> 1    3 -> 2
## intercept -2.322311 -5.269481 -1.629989 -2.175968 -4.255974 -1.192184
##
## $delta
## [1] 0.5417274 0.3293862 0.1288865
##
## $gamma
##           [,1]      [,2]      [,3]
## [1,] 0.90645971 0.08887539 0.004664898
## [2,] 0.14963129 0.76369109 0.086677616
## [3,] 0.01076032 0.23036280 0.758876882
mod4$mle

## $stepPar
##           state 1  state 2  state 3  state 4
## mean 10.096558 37.06309 129.2716 366.8990
## sd    7.071197 22.52219 67.4132 194.7513
##
## $anglePar
## NULL
##
## $beta
##           1 -> 2    1 -> 3    1 -> 4    2 -> 1    2 -> 3    2 -> 4    3 -> 1
## intercept -1.104223 -2.603545 -4.82448 -1.654802 -2.08783 -4.423115 -5.911706
##           3 -> 2    3 -> 4    4 -> 1    4 -> 2    4 -> 3
## intercept -1.42912 -2.186579 -15.37658 -2.398733 -1.236545
##
## $delta
## [1] 0.20863821 0.42001865 0.27168173 0.09966142
##
## $gamma
##           [,1]      [,2]      [,3]      [,4]
## [1,] 7.074587e-01 0.2345003 0.05235956 0.005681407
## [2,] 1.440227e-01 0.7535326 0.09340479 0.009039980
## [3,] 1.998900e-03 0.1768289 0.73826485 0.082907373
## [4,] 1.519755e-07 0.0657629 0.21023882 0.723998132



```


When we observe the states, especially in the 4-state setting, it is apparent that the states rather "mop up" the structure.

Now, considering this set of candidate models for the data set we have, how could we select the best one? What criteria to utilize besides the information criteria? Let's plot the marginal distribution of the states to see how well they fit to empirical distributions:

```
plot(mod2, breaks=60)
```


```
## Decoding states sequence... DONE
```







```
plot(mod3, breaks=60)
```

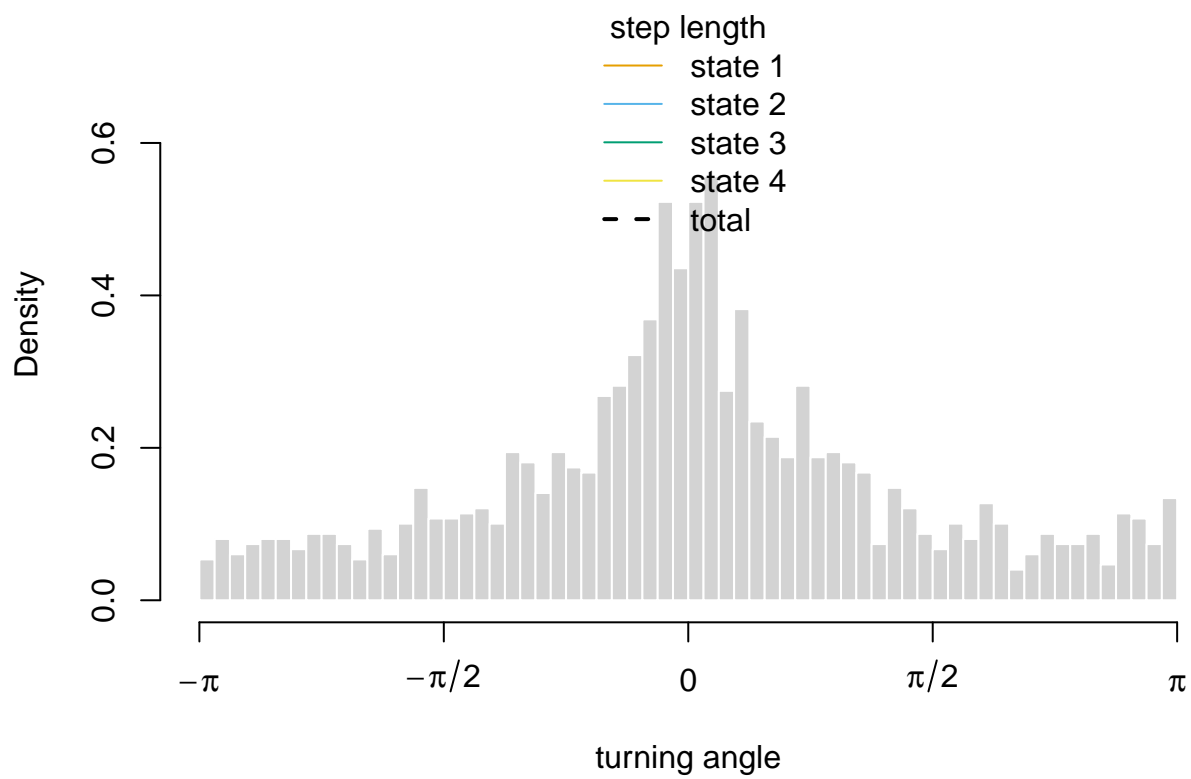

```
## Decoding states sequence... DONE
```

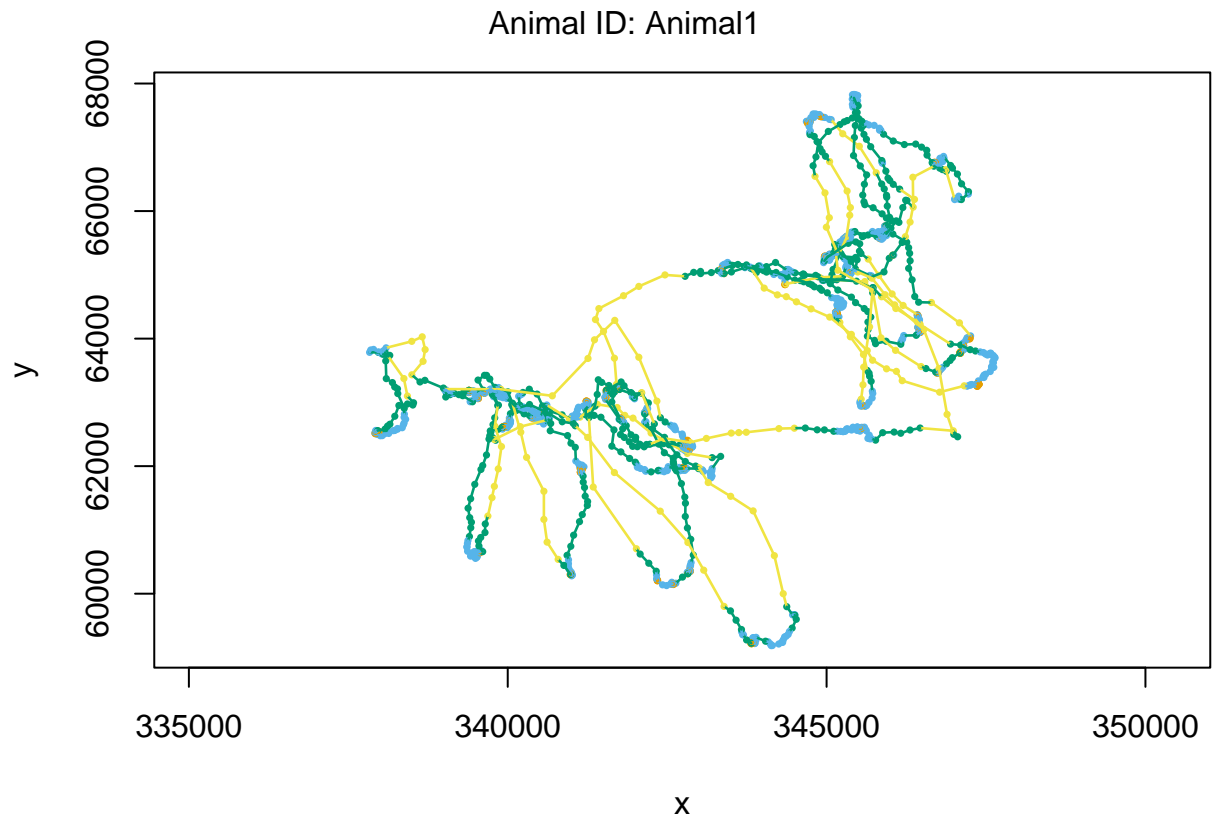




```
plot(mod4, breaks=60)
```


```
## Decoding states sequence... DONE
```





```
mu = as.vector(mod3$mle$stepPar[1,])
sigma = as.vector(mod3$mle$stepPar[2,])
delta = as.vector(mod3$mle$delta)
par(mfrow=c(1,1))
hist(data$step,probability=TRUE,breaks=60,col="light grey",xlab="step length in metres",main="Marginal o
z <- seq(0,1200,by=0.1)
lines(z,
      (delta[1]*dgamma(z,shape=mu[1]^2/sigma[1]^2,scale=sigma[1]^2/mu[1])
      + delta[2]*dgamma(z,shape=mu[2]^2/sigma[2]^2,scale=sigma[2]^2/mu[2])
      + delta[3]*dgamma(z,shape=mu[3]^2/sigma[3]^2,scale=sigma[3]^2/mu[3])),
      col='red',lwd=2)
```

Marginal distribution vs. empirical distribution



Now we move on to the residual checking for the model fit:

```
pr2 <- pseudoRes(mod2)
pr3 <- pseudoRes(mod3)
pr4 <- pseudoRes(mod4)

# Plot the residuals for the first series (e.g., pr2)
plot(pr2$stepRes, ylab = "pseudo-residuals", main = "Pseudo-Residuals for HMMs",
     col = "blue", pch = 20)


# Add a horizontal line at y=0
abline(h = 0, lty = 2, col = "grey")

# Add the residuals for the second series (pr3) in a different color
points(pr3$stepRes, col = "red", pch = 20)

# Add the residuals for the third series (pr4)
points(pr4$stepRes, col = "green", pch = 20)

# Add a legend to distinguish the series
legend("topleft", legend = c("2 states", "3 states", "4 states"),
     col = c("blue", "red", "green"), pch = 20, title = "HMM")
```

Pseudo-Residuals for HMMs



```
# Define semi-transparent colors
blue_alpha <- rgb(0, 0, 255, 100, maxColorValue = 255)
red_alpha <- rgb(255, 0, 0, 100, maxColorValue = 255)
green_alpha <- rgb(0, 255, 0, 100, maxColorValue = 255)

# Plot the residuals for the first series (pr2) using open circles
plot(pr2$stepRes, ylab = "pseudo-residuals", main = "Pseudo-Residuals for HMMs",
     col = blue_alpha, pch = 1)

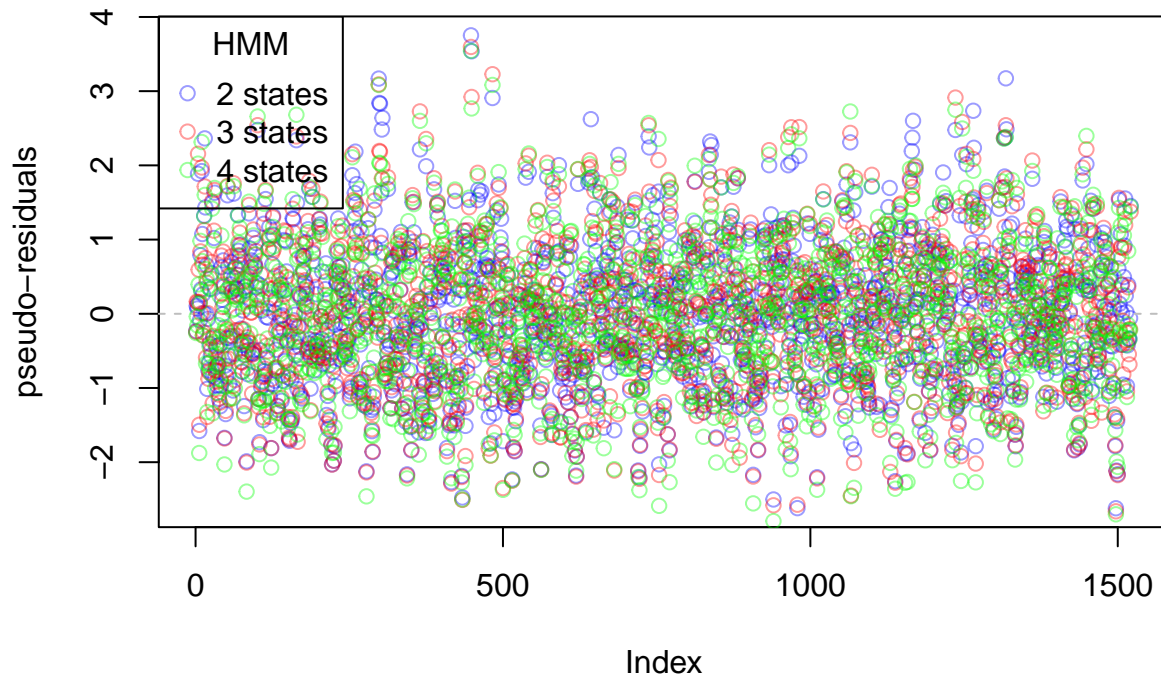
# Add a horizontal line at y=0
abline(h = 0, lty = 2, col = "grey")

# Add the residuals for the second series (pr3) with open circles
points(pr3$stepRes, col = red_alpha, pch = 1)

# Add the residuals for the third series (pr4) with open circles
points(pr4$stepRes, col = green_alpha, pch = 1)

# Add a legend with open circles
legend("topleft", legend = c("2 states", "3 states", "4 states"),
     col = c(blue_alpha, red_alpha, green_alpha), pch = 1, title = "HMM")
```

Pseudo-Residuals for HMMs



Now, after checking out these things, and that we are content with our choice of the model with 3-states. We move on to state decoding.

But there is a catch with it! There are two kinds of state decoding, namely, global and local. In local state decoding, we are asking our algorithm to point out what is the most likely state at each individual point t . And for global state decoding, we are asking the respective algorithm of ours which path is the most likely one that could occur. So instead of checking our ass at each step, we are gazing into the horizon and trying to point out which road to follow.

They can be both good and bad, depending on what we expect from them. Local decoding could possibly miss the temporal structure and as if 'overfit' to data. But it is quite precise in decoding what is going on at each point, regardless of the neighboring activity. Zig-zagging due to that is inevitable.

Meanwhile, global decoding rather smooths out the state transition process, possibly ignoring local anomalies, so as to be blind to these.

```
sp3 <- stateProbs(mod3)
round(sp3[1:50,],3)
```

```
##      [,1] [,2] [,3]
## [1,] 0.882 0.118 0.001
## [2,] 0.859 0.140 0.000
## [3,] 0.706 0.293 0.001
## [4,] 0.001 0.952 0.047
## [5,] 0.003 0.970 0.027
## [6,] 0.444 0.555 0.001
## [7,] 0.348 0.650 0.002
## [8,] 0.099 0.891 0.010
## [9,] 0.038 0.895 0.066
## [10,] 0.000 0.059 0.941
## [11,] 0.000 0.022 0.978
```

```

## [12,] 0.000 0.111 0.889
## [13,] 0.000 0.136 0.864
## [14,] 0.000 0.005 0.995
## [15,] 0.000 0.000 1.000
## [16,] 0.000 0.008 0.992
## [17,] 0.000 0.392 0.608
## [18,] 0.000 0.526 0.474
## [19,] 0.132 0.844 0.024
## [20,] 0.222 0.775 0.003
## [21,] 0.801 0.199 0.000
## [22,] 0.980 0.020 0.000
## [23,] 0.995 0.005 0.000
## [24,] 0.986 0.014 0.000
## [25,] 0.982 0.018 0.000
## [26,] 0.798 0.202 0.000
## [27,] 0.504 0.495 0.001
## [28,] 0.881 0.119 0.000
## [29,] 0.846 0.154 0.000
## [30,] 0.273 0.724 0.002
## [31,] 0.012 0.967 0.021
## [32,] 0.532 0.468 0.001
## [33,] 0.905 0.094 0.000
## [34,] 0.988 0.012 0.000
## [35,] 0.999 0.001 0.000
## [36,] 0.999 0.001 0.000
## [37,] 0.999 0.001 0.000
## [38,] 1.000 0.000 0.000
## [39,] 0.998 0.002 0.000
## [40,] 1.000 0.000 0.000
## [41,] 0.999 0.001 0.000
## [42,] 0.981 0.019 0.000
## [43,] 0.987 0.013 0.000
## [44,] 0.988 0.012 0.000
## [45,] 0.969 0.031 0.000
## [46,] 0.982 0.018 0.000
## [47,] 1.000 0.000 0.000
## [48,] 0.999 0.001 0.000
## [49,] 1.000 0.000 0.000
## [50,] 0.996 0.004 0.000


```

```

matplot(sp3[1:50,], type = "l", lty = 1, lwd = 2,
        col = c("blue", "red", "darkgreen"),
        xlab = "Time", ylab = "State Probability",
        main = "Posterior State Probabilities")
legend("topright", legend = c("State 1", "State 2", "State 3"),
      col = c("blue", "red", "darkgreen"), lty = 1, lwd = 2)

```


Posterior State Probabilities



We observe here that the probabilities are mostly gradually rising and falling, making it indeed hard to determine the most probable sequence of states.

I wonder, what difference would we observe for the first 50 values if we check the 4-state model. Let's PLOT!


```
sp4 <- stateProbs(mod4)
round(sp4[1:50,],3)
```

```
##      [,1] [,2] [,3] [,4]
## [1,] 0.027 0.955 0.018 0.000
## [2,] 0.015 0.971 0.013 0.000
## [3,] 0.027 0.896 0.077 0.000
## [4,] 0.000 0.001 0.966 0.033
## [5,] 0.000 0.011 0.966 0.023
## [6,] 0.006 0.924 0.070 0.000
## [7,] 0.003 0.898 0.099 0.000
## [8,] 0.000 0.616 0.382 0.002
## [9,] 0.000 0.452 0.521 0.026
## [10,] 0.000 0.000 0.166 0.834
## [11,] 0.000 0.000 0.116 0.884
## [12,] 0.000 0.000 0.300 0.700
## [13,] 0.000 0.000 0.319 0.681
## [14,] 0.000 0.000 0.018 0.982
## [15,] 0.000 0.000 0.000 1.000
## [16,] 0.000 0.000 0.020 0.980
## [17,] 0.000 0.001 0.593 0.407
## [18,] 0.000 0.000 0.701 0.299
## [19,] 0.000 0.647 0.347 0.006
## [20,] 0.000 0.715 0.284 0.002
## [21,] 0.017 0.957 0.026 0.000
```

```
## [22,] 0.145 0.854 0.001 0.000
## [23,] 0.153 0.846 0.001 0.000
## [24,] 0.040 0.957 0.004 0.000
## [25,] 0.139 0.859 0.002 0.000
## [26,] 0.006 0.940 0.054 0.000
## [27,] 0.000 0.733 0.265 0.002
## [28,] 0.138 0.853 0.009 0.000
## [29,] 0.133 0.850 0.017 0.000
## [30,] 0.000 0.537 0.461 0.002
## [31,] 0.000 0.029 0.949 0.023
## [32,] 0.000 0.818 0.181 0.001
## [33,] 0.077 0.914 0.009 0.000
## [34,] 0.420 0.580 0.000 0.000
## [35,] 0.872 0.128 0.000 0.000
## [36,] 0.940 0.060 0.000 0.000
## [37,] 0.956 0.044 0.000 0.000
## [38,] 0.978 0.022 0.000 0.000
## [39,] 0.932 0.068 0.000 0.000
## [40,] 0.963 0.037 0.000 0.000
## [41,] 0.827 0.173 0.000 0.000
## [42,] 0.019 0.968 0.013 0.000
## [43,] 0.014 0.981 0.004 0.000
## [44,] 0.024 0.973 0.003 0.000
## [45,] 0.002 0.988 0.010 0.000
## [46,] 0.054 0.946 0.001 0.000
## [47,] 0.940 0.060 0.000 0.000
## [48,] 0.959 0.041 0.000 0.000
## [49,] 0.981 0.019 0.000 0.000
## [50,] 0.898 0.102 0.000 0.000
```

```
matplot(sp4[1:50,], type = "l", lty = 1, lwd = 2,
        col = c("blue", "red", "darkgreen", "purple"),
        xlab = "Time", ylab = "State Probability",
        main = "Posterior State Probabilities")
legend("topright", legend = c("State 1", "State 2", "State 3", "State 4"),
        col = c("blue", "red", "darkgreen", "purple"), lty = 1, lwd = 2)
```

Posterior State Probabilities



From the visual examination, the most we can claim is that 4-state setting makes the state probabilities act more wiggly and less apparent. At least, it can be claimed that the state probabilities seem unsure and rather noisy.

```
plotStates(mod3)
```

```
## Decoding states sequence... DONE  
## Computing states probabilities... DONE
```

Animal ID: Animal1

